

```


In [1]: import torch
import torchvision
from torch.utils.data import DataLoader
from torchvision.datasets import MNIST
from torchvision import transforms
from torch import nn
import math
import os
os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"

In [2]: # batch_size 超参, 根据硬件配置相应大小
batch_size = 32 #批处理大小
trans_img = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

In [3]: # MNIST 数据集每张图片是灰度图片, 大小为 28x28
trainset = MNIST('data', train=True, download=True, transform=trans_img)
testset = MNIST('data', train=False, download=True, transform=trans_img)
train_loader = DataLoader(trainset, batch_size=batch_size,
    shuffle=True, num_workers=1)
test_loader = DataLoader(testset, batch_size=batch_size,
    shuffle=True, num_workers=1) # 运行环境

In [4]: import matplotlib.pyplot as plt
plt.ion()
cnt = 0
for (img_batch, label) in train_loader:
    cnt += 1
    if cnt > 10:
        break
    fig, ax = plt.subplots(
        nrows=4,
        ncols=8,
        sharex=True,
        sharey=True, )
    ax = ax.flatten()
    for i in range(32):
        img = img_batch[i].numpy().reshape(28, 28)
        ax[i].imshow(img, cmap='Greys', interpolation='nearest')
    ax[0].set_xticks([])
    ax[0].set_yticks([])
    plt.show()
    plt.close()
plt.ioff()

```



```
In [5]: class Net(nn.Module):
        def __init__(self):
            super(Net, self).__init__()

            self.features = nn.Sequential(
                nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
                nn.BatchNorm2d(32),
                nn.ReLU(inplace=True), #inplace = True ,
                nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1),
                nn.BatchNorm2d(32),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(kernel_size=2, stride=2),
                nn.Conv2d(32, 64, kernel_size=3, padding=1),
                nn.BatchNorm2d(64),
                nn.ReLU(inplace=True),
                nn.Conv2d(64, 64, kernel_size=3, padding=1),
                nn.BatchNorm2d(64),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(kernel_size=2, stride=2)
            )
            self.classifier = nn.Sequential(
                nn.Dropout(p = 0.5),
                nn.Linear(64 * 7 * 7, 512),
                nn.BatchNorm1d(512),
                nn.ReLU(inplace=True),
                nn.Dropout(p = 0.5),
                nn.Linear(512, 512),
                nn.BatchNorm1d(512),
                nn.ReLU(inplace=True),
                nn.Dropout(p = 0.5),
                nn.Linear(512, 10),
            )
```

```
for m in self.features.children():
    if isinstance(m, nn.Conv2d):
        n = m.kernel_size[0] * m.kernel_size[1] * m.out_channels
        m.weight.data.normal_(0, math.sqrt(2. / n))
    elif isinstance(m, nn.BatchNorm2d):
        m.weight.data.fill_(1)
        m.bias.data.zero_()
for m in self.classifier.children():
    if isinstance(m, nn.Linear):
        nn.init.xavier_uniform(m.weight)
    elif isinstance(m, nn.BatchNorm1d):
        m.weight.data.fill_(1)
        m.bias.data.zero_()
def forward(self, x):
    x = self.features(x)
    x = x.view(x.size(0), -1) #这句话是说将最后一次卷积的输出拉伸为一行
    x = self.classifier(x)
    return x
```

```

In [6]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        from torch import optim #优化器
        from torch.autograd import Variable
        import warnings
        warnings.filterwarnings("ignore", category=UserWarning)
        model = Net().to(device) #GPU
        learning_rate = 0.001
        criterion = nn.CrossEntropyLoss(size_average=False)
        optimizer = optim.SGD(model.parameters(), lr = learning_rate)
        # 总的训练轮数
        epochs = 50
        train_losses=[]
        test_losses=[]
        for epoch in range(epochs):
            running_loss, running_acc = 0., 0.
            for (img, label) in train_loader:
                img = Variable(img).to(device)
                label = Variable(label).to(device)

                optimizer.zero_grad() #梯度归零;

                output = model(img) ##前向传播
                loss = criterion(output, label) #计算 loss

                loss.backward() #反向传播, 计算当前梯度;
                optimizer.step() #反向传播, 计算当前梯度;
                running_loss += loss.item() #item() 取出张量具体位置的元素元素值
                _, predict = torch.max(output, 1)
                #_, predicted = torch.max(outputs.data, dim): 返回最大值所在索引, dim=1 时, 按行返回最大值所在索引
                correct_num = (predict == label).sum()
                running_acc += correct_num.item()

```

```

        running_loss /= len(trainset)
        running_acc /= len(trainset)

        with torch.no_grad(): #不求梯度
            test_loss, test_acc = 0., 0.
            for images, labels in test_loader:
                images = Variable(images).to(device)
                labels = Variable(labels).to(device)
                output = model(images)
                loss = criterion(output, labels)
                test_loss += loss.item()
                _, predict = torch.max(output, 1)
                correct_num = (predict == labels).sum()
                test_acc += correct_num.item()

            test_loss /= len(testset)
            test_acc /= len(testset)

            train_losses.append(running_loss)
            test_losses.append(test_loss)
            print("Epoch: {}/{}\n".format(epoch+1, epochs),
                  "Training Loss: {:.3f}\n".format(train_losses[-1]),
                  "Training Accuracy: {:.3f} %\n".format(100*running_acc),
                  "Test Loss: {:.3f}\n".format(test_losses[-1]),
                  "Test Accuracy: {:.3f} %\n".format(100*test_acc))
        # 保存模型
        torch.save(model, 'conv.pth.tar')

```

```

Epoch: 1/50.. Training Loss: 0.294.. Training Accuracy: 90.587 % Test Loss: 0.131.. Test Accuracy: 95.920 %
Epoch: 2/50.. Training Loss: 0.119.. Training Accuracy: 96.377 % Test Loss: 0.096.. Test Accuracy: 96.920 %
Epoch: 3/50.. Training Loss: 0.096.. Training Accuracy: 97.008 % Test Loss: 0.081.. Test Accuracy: 97.530 %
Epoch: 4/50.. Training Loss: 0.081.. Training Accuracy: 97.527 % Test Loss: 0.078.. Test Accuracy: 97.540 %
Epoch: 5/50.. Training Loss: 0.071.. Training Accuracy: 97.815 % Test Loss: 0.064.. Test Accuracy: 97.950 %
Epoch: 6/50.. Training Loss: 0.063.. Training Accuracy: 98.033 % Test Loss: 0.059.. Test Accuracy: 98.080 %
Epoch: 7/50.. Training Loss: 0.058.. Training Accuracy: 98.212 % Test Loss: 0.058.. Test Accuracy: 98.200 %
Epoch: 8/50.. Training Loss: 0.055.. Training Accuracy: 98.318 % Test Loss: 0.050.. Test Accuracy: 98.440 %
Epoch: 9/50.. Training Loss: 0.051.. Training Accuracy: 98.450 % Test Loss: 0.050.. Test Accuracy: 98.280 %
Epoch: 10/50.. Training Loss: 0.046.. Training Accuracy: 98.548 % Test Loss: 0.049.. Test Accuracy: 98.460 %
Epoch: 11/50.. Training Loss: 0.045.. Training Accuracy: 98.655 % Test Loss: 0.048.. Test Accuracy: 98.440 %
Epoch: 12/50.. Training Loss: 0.046.. Training Accuracy: 98.582 % Test Loss: 0.049.. Test Accuracy: 98.590 %
Epoch: 13/50.. Training Loss: 0.042.. Training Accuracy: 98.687 % Test Loss: 0.042.. Test Accuracy: 98.710 %
Epoch: 14/50.. Training Loss: 0.038.. Training Accuracy: 98.802 % Test Loss: 0.040.. Test Accuracy: 98.790 %
Epoch: 15/50.. Training Loss: 0.038.. Training Accuracy: 98.823 % Test Loss: 0.042.. Test Accuracy: 98.750 %
Epoch: 16/50.. Training Loss: 0.034.. Training Accuracy: 98.923 % Test Loss: 0.040.. Test Accuracy: 98.860 %
Epoch: 17/50.. Training Loss: 0.035.. Training Accuracy: 98.962 % Test Loss: 0.037.. Test Accuracy: 98.770 %
Epoch: 18/50.. Training Loss: 0.033.. Training Accuracy: 99.000 % Test Loss: 0.035.. Test Accuracy: 98.870 %
Epoch: 19/50.. Training Loss: 0.033.. Training Accuracy: 98.953 % Test Loss: 0.038.. Test Accuracy: 98.840 %
Epoch: 20/50.. Training Loss: 0.029.. Training Accuracy: 99.068 % Test Loss: 0.035.. Test Accuracy: 98.870 %
Epoch: 21/50.. Training Loss: 0.031.. Training Accuracy: 99.077 % Test Loss: 0.038.. Test Accuracy: 98.800 %
Epoch: 22/50.. Training Loss: 0.028.. Training Accuracy: 99.103 % Test Loss: 0.037.. Test Accuracy: 98.900 %
Epoch: 23/50.. Training Loss: 0.027.. Training Accuracy: 99.168 % Test Loss: 0.033.. Test Accuracy: 98.980 %
Epoch: 24/50.. Training Loss: 0.027.. Training Accuracy: 99.152 % Test Loss: 0.034.. Test Accuracy: 98.970 %
Epoch: 25/50.. Training Loss: 0.026.. Training Accuracy: 99.187 % Test Loss: 0.037.. Test Accuracy: 98.850 %
Epoch: 26/50.. Training Loss: 0.026.. Training Accuracy: 99.210 % Test Loss: 0.034.. Test Accuracy: 99.010 %
Epoch: 27/50.. Training Loss: 0.026.. Training Accuracy: 99.157 % Test Loss: 0.035.. Test Accuracy: 98.870 %
Epoch: 28/50.. Training Loss: 0.024.. Training Accuracy: 99.250 % Test Loss: 0.030.. Test Accuracy: 99.040 %
Epoch: 29/50.. Training Loss: 0.025.. Training Accuracy: 99.218 % Test Loss: 0.034.. Test Accuracy: 98.960 %
Epoch: 30/50.. Training Loss: 0.022.. Training Accuracy: 99.315 % Test Loss: 0.037.. Test Accuracy: 98.840 %
Epoch: 31/50.. Training Loss: 0.022.. Training Accuracy: 99.253 % Test Loss: 0.027.. Test Accuracy: 99.120 %


```

```

In [7]: plt.plot(train_losses, label='Training loss')
plt.plot(test_losses, label='Validation loss')
plt.xlabel("epochs")
plt.ylabel("loss")
plt.legend(frameon=False)

```

Out[7]: <matplotlib.legend.Legend at 0x23ae6452c70>



```

In [9]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = torch.load('conv.pth.tar')
print('testing cnn model')
testloss, testacc = 0., 0.
for (img, label) in test_loader:
    img = Variable(img).to(device)
    label = Variable(label).to(device)
    out = model(img)
    loss = criterion(out, label)
    testloss += loss.item()
    _, predict = torch.max(out, 1)
    correct_num = (predict == label).sum()
    testacc += correct_num.item()
testloss /= len(testset)
testacc /= len(testset)
print('cnn model, Test: Loss: %.5f, Acc: %.2f' %
      (testloss, 100 * testacc))

```

```

testing cnn model
cnn model, Test: Loss: 0.03346, Acc: 99.09

```